

ML - Obesity Prediction

May 23, 2024

1 Importing Packages

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

2 Loading Data & Basic Analysis

```
[2]: data= pd.read_csv('ObesityDataSet_raw_and_data_sinthetic.csv')
df = data.copy()
df.head()
```

```
[2]:
```

	Age	Gender	Height	Weight	CALC	FAVC	FCVC	NCP	SCC	SMOKE	CH2O	\
0	21.0	Female	1.62	64.0	no	no	2.0	3.0	no	no	2.0	
1	21.0	Female	1.52	56.0	Sometimes	no	3.0	3.0	yes	yes	3.0	
2	23.0	Male	1.80	77.0	Frequently	no	2.0	3.0	no	no	2.0	
3	27.0	Male	1.80	87.0	Frequently	no	3.0	3.0	no	no	2.0	
4	22.0	Male	1.78	89.8	Sometimes	no	2.0	1.0	no	no	2.0	

	family_history_with_overweight	FAF	TUE	CAEC	MTRANS	\
0	yes	0.0	1.0	Sometimes	Public_Transportation	
1	yes	3.0	0.0	Sometimes	Public_Transportation	
2	yes	2.0	1.0	Sometimes	Public_Transportation	
3	no	2.0	0.0	Sometimes	Walking	
4	no	0.0	0.0	Sometimes	Public_Transportation	

	NObeyesdad
0	Normal_Weight
1	Normal_Weight
2	Normal_Weight
3	Overweight_Level_I
4	Overweight_Level_II

```
[3]: df.tail()
```

```
[3]:
```

	Age	Gender	Height	Weight	CALC	FAVC	FCVC	NCP	SCC	\
2106	20.976842	Female	1.710730	131.408528	Sometimes	yes	3.0	3.0	no	
2107	21.982942	Female	1.748584	133.742943	Sometimes	yes	3.0	3.0	no	
2108	22.524036	Female	1.752206	133.689352	Sometimes	yes	3.0	3.0	no	
2109	24.361936	Female	1.739450	133.346641	Sometimes	yes	3.0	3.0	no	
2110	23.664709	Female	1.738836	133.472641	Sometimes	yes	3.0	3.0	no	

	SMOKE	CH20	family_history_with_overweight	FAF	TUE	\
2106	no	1.728139		yes	1.676269	0.906247
2107	no	2.005130		yes	1.341390	0.599270
2108	no	2.054193		yes	1.414209	0.646288
2109	no	2.852339		yes	1.139107	0.586035
2110	no	2.863513		yes	1.026452	0.714137

	CAEC	MTRANS	NObeyesdad
2106	Sometimes	Public_Transportation	Obesity_Type_III
2107	Sometimes	Public_Transportation	Obesity_Type_III
2108	Sometimes	Public_Transportation	Obesity_Type_III
2109	Sometimes	Public_Transportation	Obesity_Type_III
2110	Sometimes	Public_Transportation	Obesity_Type_III

```
[4]: df.shape
```

```
[4]: (2111, 17)
```

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    2111 non-null   float64
1   Gender                                2111 non-null   object
2   Height                                2111 non-null   float64
3   Weight                                2111 non-null   float64
4   CALC                                  2111 non-null   object
5   FAVC                                  2111 non-null   object
6   FCVC                                  2111 non-null   float64
7   NCP                                    2111 non-null   float64
8   SCC                                    2111 non-null   object
9   SMOKE                                  2111 non-null   object
10  CH20                                   2111 non-null   float64
11  family_history_with_overweight        2111 non-null   object
12  FAF                                    2111 non-null   float64
13  TUE                                    2111 non-null   float64
14  CAEC                                  2111 non-null   object
15  MTRANS                                2111 non-null   object
```

```

16 NObeyesdad                2111 non-null    object
dtypes: float64(8), object(9)
memory usage: 280.5+ KB

```

```
[6]: df.describe().T
```

```

[6]:      count      mean      std   min    25%    50%    75%  \
Age      2111.0   24.312600   6.345968  14.00  19.947192  22.777890  26.000000
Height    2111.0    1.701677   0.093305   1.45   1.630000   1.700499   1.768464
Weight    2111.0   86.586058  26.191172  39.00  65.473343  83.000000  107.430682
FCVC      2111.0    2.419043   0.533927   1.00   2.000000   2.385502   3.000000
NCP       2111.0    2.685628   0.778039   1.00   2.658738   3.000000   3.000000
CH20      2111.0    2.008011   0.612953   1.00   1.584812   2.000000   2.477420
FAF       2111.0    1.010298   0.850592   0.00   0.124505   1.000000   1.666678
TUE       2111.0    0.657866   0.608927   0.00   0.000000   0.625350   1.000000

      max
Age      61.00
Height    1.98
Weight   173.00
FCVC      3.00
NCP       4.00
CH20      3.00
FAF       3.00
TUE       2.00

```

```
[7]: df.describe(include=["O"]).T
```

```

[7]:      count  unique      top  freq
Gender      2111      2      Male  1068
CALC        2111      4  Sometimes  1401
FAVC        2111      2      yes   1866
SCC          2111      2      no   2015
SMOKE        2111      2      no   2067
family_history_with_overweight  2111      2      yes   1726
CAEC         2111      4  Sometimes  1765
MTRANS       2111      5  Public_Transportation  1580
NObeyesdad   2111      7  Obesity_Type_I      351

```

```
[8]: df.isnull().sum()
```

```

[8]: Age      0
Gender      0
Height      0
Weight      0
CALC        0
FAVC        0
FCVC        0

```

```

NCP                                0
SCC                                0
SMOKE                              0
CH20                               0
family_history_with_overweight    0
FAF                                0
TUE                                0
CAEC                               0
MTRANS                             0
NObeyesdad                         0
dtype: int64

```

```
[9]: df.duplicated().sum()
```

```
[9]: 24
```

- Keeping duplicates values as we don't have a primary key and the rows could be of different patients

3 Exploratory Data Analysis

```
[10]: df.columns
```

```
[10]: Index(['Age', 'Gender', 'Height', 'Weight', 'CALC', 'FAVC', 'FCVC', 'NCP',
          'SCC', 'SMOKE', 'CH20', 'family_history_with_overweight', 'FAF', 'TUE',
          'CAEC', 'MTRANS', 'NObeyesdad'],
          dtype='object')
```


3.1 Correlation matrix

```
[11]: # Select only numeric columns for correlation analysis
numerical_data = df.select_dtypes(include=['float64'])

# Calculate the correlation matrix
correlation_matrix = numerical_data.corr()
```

```
[12]: plt.figure(figsize=(8, 12))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwarm',
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

plt.title('Heatmap of Correlation Matrix')
plt.show()
```




[]:

3.2 Height and Weight Distribution


```
[13]: plt.hist(df['Height'], bins=20, edgecolor='black')
plt.title('Distribution of Height')
plt.xlabel('Height')
plt.ylabel('Frequency')

plt.show()
```




```
[14]: plt.hist(df['Weight'], bins=20, edgecolor='black')
plt.title('Distribution of Weight')
plt.xlabel('Weight')
plt.ylabel('Frequency')

plt.show()
```



3.3 Top 10 Ages with Highest Weight

```
[15]: Top10_age = df.groupby('Age')['Weight'].mean().sort_values(ascending=False).  
      ↪head(10)  
  
plt.figure(figsize=(10, 6))  
Top10_age.plot(kind='bar', color='skyblue')  
  
plt.xticks(rotation=45)  
plt.show()  
Top10_age
```



```
[15]: Age
      21.900120    165.057269
      19.472190    160.935351
      21.521294    160.639405
      20.989016    155.872093
      21.291969    155.242672
      20.327723    154.618446
      20.102241    153.959945
      20.438478    153.149491
      18.112503    152.720545
      18.120739    152.567671
      Name: Weight, dtype: float64
```

```
[ ]:
```

3.4 Distribution of CAEC values

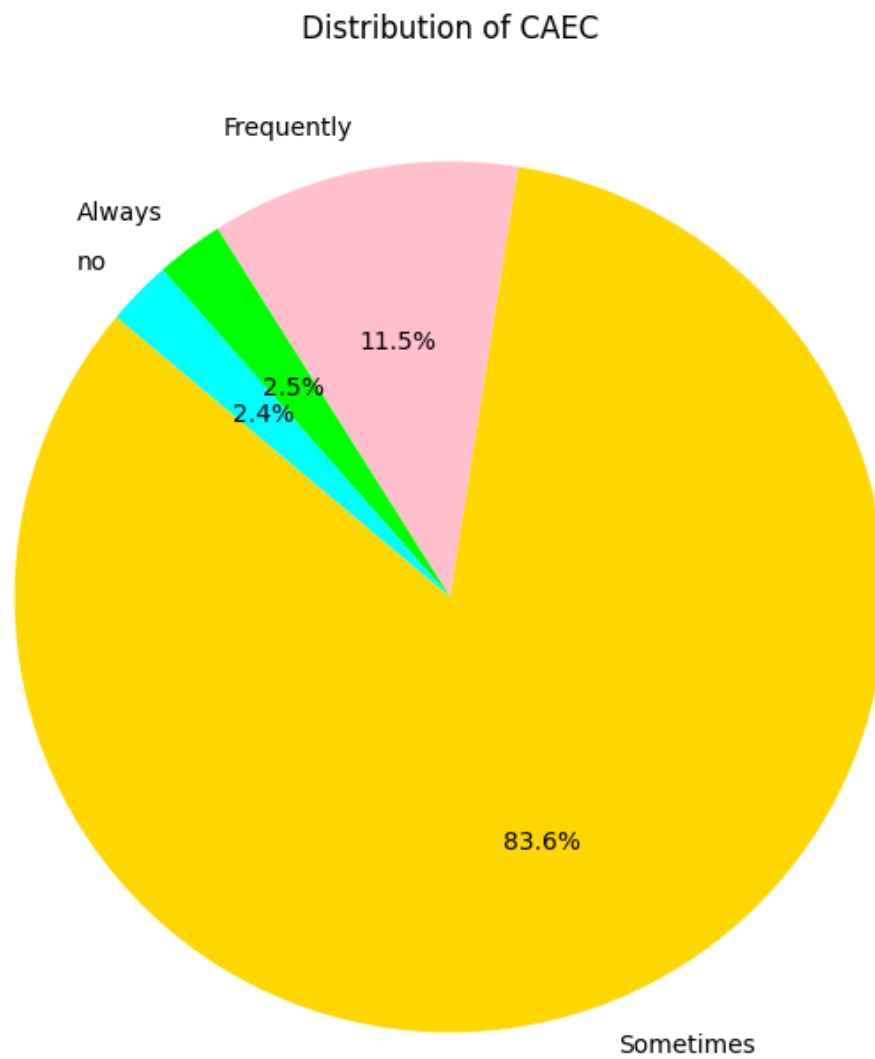
```
[16]: caec_count = df['CAEC'].value_counts()

plt.figure(figsize=(8, 8))
caec_count.plot(kind='pie', autopct='%1.1f%%', startangle=140, colors=['gold', 'pink', 'lime', 'cyan'])
```



```
plt.title('Distribution of CAEC')
plt.ylabel('')
plt.show()

caec_count
```




```
[16]: CAEC
      Sometimes    1765
      Frequently    242
      Always        53
      no            51
```

Name: count, dtype: int64

[]:

3.5 Average Ages with Family History with Overweight

```
[17]: avg_age_with_familyHistory = df.  
      ↳groupby('family_history_with_overweight')['Age'].mean()  
  
plt.figure(figsize=(10, 6))  
avg_age_with_familyHistory.plot(kind='bar', color='skyblue')  
plt.title('Average Ages with Family History with Overweight')  
plt.xlabel('Family History with Overweight')  
plt.ylabel('Average Age')  
plt.xticks(rotation=0)  
  
plt.show()  
avg_age_with_familyHistory
```



```
[17]: family_history_with_overweight  
no    21.549015  
yes    24.929043  
Name: Age, dtype: float64
```

```
[ ]:
```

4 Data Preprocessing

```
[18]: from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder
      from sklearn.model_selection import train_test_split
```

```
[19]: categorical_cols = df.select_dtypes(include=['object']).columns.tolist()
      categorical_cols
```

```
[19]: ['Gender',
      'CALC',
      'FAVC',
      'SCC',
      'SMOKE',
      'family_history_with_overweight',
      'CAEC',
      'MTRANS',
      'NObeyesdad']
```

```
[20]: continuous_cols = df.select_dtypes(include=['float64']).columns.tolist()
      continuous_cols
```

```
[20]: ['Age', 'Height', 'Weight', 'FCVC', 'NCP', 'CH20', 'FAF', 'TUE']
```

```
[21]: # Applying Label Encoding to categorical columns
      label_encoder = LabelEncoder()

      for col in categorical_cols:
          df[col] = label_encoder.fit_transform(df[col])
```

```
[22]: # Applying Standard Scaling to continuous columns
      scaler = StandardScaler()
      df[continuous_cols] = scaler.fit_transform(df[continuous_cols])
```

```
[23]: # Applying One-Hot Encoding to nominal categorical variable with more than two
      ↪ categories.
      onehot_encoder = OneHotEncoder(categories='auto', sparse_output=False)

      nominal_cols = ['CALC', 'CAEC', 'MTRANS']
      df1 = pd.get_dummies(df, columns=nominal_cols)
```

```
[24]: df1
```

```
[24]:
```

	Age	Gender	Height	Weight	FAVC	FCVC	NCP	SCC	\
0	-0.522124	0	-0.875589	-0.862558	0	-0.785019	0.404153	0	
1	-0.522124	0	-1.947599	-1.168077	0	1.088342	0.404153	1	

2	-0.206889	1	1.054029	-0.366090	0	-0.785019	0.404153	0
3	0.423582	1	1.054029	0.015808	0	1.088342	0.404153	0
4	-0.364507	1	0.839627	0.122740	0	-0.785019	-2.167023	0
...
2106	-0.525774	0	0.097045	1.711763	1	1.088342	0.404153	0
2107	-0.367195	0	0.502844	1.800914	1	1.088342	0.404153	0
2108	-0.281909	0	0.541672	1.798868	1	1.088342	0.404153	0
2109	0.007776	0	0.404927	1.785780	1	1.088342	0.404153	0
2110	-0.102119	0	0.398344	1.790592	1	1.088342	0.404153	0

	SMOKE	CH2O	...	CALC_3	CAEC_0	CAEC_1	CAEC_2	CAEC_3	MTRANS_0	\
0	0	-0.013073	...	True	False	False	True	False	False	
1	1	1.618759	...	False	False	False	True	False	False	
2	0	-0.013073	...	False	False	False	True	False	False	
3	0	-0.013073	...	False	False	False	True	False	False	
4	0	-0.013073	...	False	False	False	True	False	False	
...	
2106	0	-0.456705	...	False	False	False	True	False	False	
2107	0	-0.004702	...	False	False	False	True	False	False	
2108	0	0.075361	...	False	False	False	True	False	False	
2109	0	1.377801	...	False	False	False	True	False	False	
2110	0	1.396035	...	False	False	False	True	False	False	

	MTRANS_1	MTRANS_2	MTRANS_3	MTRANS_4
0	False	False	True	False
1	False	False	True	False
2	False	False	True	False
3	False	False	False	True
4	False	False	True	False
...
2106	False	False	True	False
2107	False	False	True	False
2108	False	False	True	False
2109	False	False	True	False
2110	False	False	True	False

[2111 rows x 27 columns]

```
[25]: # Split the dataset into features and target variable
X = df1.drop('NObeyesdad', axis=1)
y = df1['NObeyesdad']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

(1688, 26) (423, 26) (1688,) (423,)

[]:

5 Algorithm Search

```
[26]: from sklearn.linear_model import LogisticRegression
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.svm import SVC
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.metrics import classification_report, accuracy_score, \
          confusion_matrix
      from sklearn.preprocessing import LabelEncoder
```

```
[27]: classifiers = {
      "Logistic Regression": LogisticRegression(max_iter=200),
      "Decision Tree": DecisionTreeClassifier(),
      "Random Forest": RandomForestClassifier(),
      "Support Vector Machine": SVC(),
      "K-Nearest Neighbors": KNeighborsClassifier()
      }

results = {}
for name, clf in classifiers.items():
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    target_names = label_encoder.classes_.astype(str)
    report = classification_report(y_test, y_pred, target_names=target_names)
    results[name] = {
        "accuracy": accuracy,
        "classification_report": report
    }
    print(f"Classifier: {name}\nAccuracy: {accuracy}\n")
    print(f"Classification Report:\n{report}\n")
```

Classifier: Logistic Regression

Accuracy: 0.8699763593380615

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.84	1.00	0.91	56
Normal_Weight	0.91	0.63	0.74	62
Obesity_Type_I	0.93	0.90	0.92	78
Obesity_Type_II	0.90	0.97	0.93	58

Obesity_Type_III	1.00	1.00	1.00	63
Overweight_Level_I	0.75	0.75	0.75	56
Overweight_Level_II	0.74	0.84	0.79	50
accuracy			0.87	423
macro avg	0.87	0.87	0.86	423
weighted avg	0.87	0.87	0.87	423

Classifier: Decision Tree
Accuracy: 0.9456264775413712

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.93	0.98	0.96	56
Normal_Weight	0.90	0.90	0.90	62
Obesity_Type_I	0.95	0.94	0.94	78
Obesity_Type_II	0.95	0.95	0.95	58
Obesity_Type_III	1.00	1.00	1.00	63
Overweight_Level_I	0.93	0.91	0.92	56
Overweight_Level_II	0.96	0.94	0.95	50
accuracy			0.95	423
macro avg	0.95	0.95	0.95	423
weighted avg	0.95	0.95	0.95	423

Classifier: Random Forest
Accuracy: 0.950354609929078

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.98	0.96	0.97	56
Normal_Weight	0.87	0.94	0.90	62
Obesity_Type_I	0.99	0.95	0.97	78
Obesity_Type_II	0.98	0.98	0.98	58
Obesity_Type_III	1.00	1.00	1.00	63
Overweight_Level_I	0.88	0.88	0.88	56
Overweight_Level_II	0.96	0.94	0.95	50
accuracy			0.95	423
macro avg	0.95	0.95	0.95	423
weighted avg	0.95	0.95	0.95	423

Classifier: Support Vector Machine

Accuracy: 0.933806146572104

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.96	0.96	0.96	56
Normal_Weight	0.84	0.87	0.86	62
Obesity_Type_I	0.96	0.97	0.97	78
Obesity_Type_II	0.97	0.98	0.97	58
Obesity_Type_III	1.00	1.00	1.00	63
Overweight_Level_I	0.84	0.84	0.84	56
Overweight_Level_II	0.96	0.88	0.92	50
accuracy			0.93	423
macro avg	0.93	0.93	0.93	423
weighted avg	0.93	0.93	0.93	423

Classifier: K-Nearest Neighbors

Accuracy: 0.8321513002364066

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.78	0.95	0.85	56
Normal_Weight	0.85	0.37	0.52	62
Obesity_Type_I	0.84	0.94	0.88	78
Obesity_Type_II	0.90	0.97	0.93	58
Obesity_Type_III	0.98	1.00	0.99	63
Overweight_Level_I	0.74	0.82	0.78	56
Overweight_Level_II	0.72	0.76	0.74	50
accuracy			0.83	423
macro avg	0.83	0.83	0.81	423
weighted avg	0.84	0.83	0.82	423

[]:

6 Best Algorithm

- Random Forest has the Best Accuracy: 95%

```
[29]: model = RandomForestClassifier(random_state=42)
      model.fit(X_train, y_train)
```

```

y_pred_ = model.predict(X_test)

myAccuracy= accuracy_score(y_test, y_pred_)

confusionReport= classification_report(y_test, y_pred_)

confusionMatrix= confusion_matrix(y_test, y_pred_)

print('Random Forest')
print('Classification Report')
print(confusionReport)
print(f'Accuracy: {myAccuracy}')
print('Confusion Matrix')
print(confusionMatrix)

```

Random Forest

Classification Report

	precision	recall	f1-score	support
0	1.00	0.96	0.98	56
1	0.88	0.95	0.91	62
2	0.99	0.94	0.96	78
3	0.97	0.98	0.97	58
4	1.00	1.00	1.00	63
5	0.94	0.89	0.92	56
6	0.91	0.96	0.93	50
accuracy			0.96	423
macro avg	0.95	0.96	0.95	423
weighted avg	0.96	0.96	0.96	423

Accuracy: 0.9550827423167849

Confusion Matrix

```

[[54  2  0  0  0  0  0]
 [ 0 59  0  0  0  1  2]
 [ 0  1 73  2  0  0  2]
 [ 0  0  1 57  0  0  0]
 [ 0  0  0  0 63  0  0]
 [ 0  5  0  0  0 50  1]
 [ 0  0  0  0  0  2 48]]

```

[]: